

animehacker TQ3_0 Final Test Report

Document control	Value
Route ID	animehacker-tq3-0
Revision	1.1
Generated date	2026-08-31
Evidence IDs	animehacker-3409c4fff1f57d4b988a, animehacker-ca0c684bb478b8527653, animehacker-34ab10f0a8210699f64b, animehacker-d0a6c059ef53abfb592a, animehacker-79e759f5588cae2c50d5, animehacker-e73b2f1205e6c7f8d4d3, animehacker-9088849e3d9565254d7d, animehacker-44a8d70b058f1c5cac52, animehacker-6517f74c71a0537eb162, animehacker-c1863347fee0fae7618f, animehacker-17b6142958ab2916e9ed, animehacker-d25f303e05b5147dc9ca, animehacker-4ee219bf5ee960cb094d

1. Title and document control

This publication is controlled by WB-03 revision 1.5 and the authenticated source set listed in the Evidence section.

DC-01 — Document control

Field	Value
Route	animehacker-tq3-0
Campaign	wb-03-v1.5-2026-07-18
Controlled source	WB-03 v1.5
Canonical format	Markdown

2. Technical summary

The final authority records seven completed runnable configurations and three resolved safety classifications. Zero unresolved failures does not mean zero historical failure attempts: eleven resolved source-ledger events and rejected runtime evidence remain auditable.

3. Key findings and decision-relevant evidence

The key findings below separate completed runtime evidence, safety classifications, and rejected historical evidence.

KF-01 — Decision-relevant findings

Finding	Evidence-bound conclusion
Runtime	Seven formal rows have exactly three explicitly included samples each.
Safety	AH-06, AH-07, and AH-10 remain safety-classified; no request metrics are fabricated.
TQ3_0	CPU activation is validated; AH-09 is CPU-resident TQ3 KV with 1/41 SYCL layers offloaded.
Rejected evidence	Superseded AH-09 summaries are excluded from statistics but retained in the attempt/failure ledger.

4. Repository, branch, commit, build, hardware, and software identity

The campaign identity is fixed to the recorded repository, machine, and operating-system environment.

SYS-01 — Controlled identity

Field	Value
Repository	https://github.com/animehacker/llama-turboquant
Commit	5bc5ed3bdc25003aa9f07422753a7b8d4f9190fc
CPU	12th Gen Intel Core i5-12450H
RAM	16.0 GB installed; 15.7 GB usable
GPU	Intel UHD Graphics; shared system memory
OS	Windows 11 Home 10.0.26200 build 26200

5. Objectives, scope, test matrix, and execution sequence

WB-03 defines AH-01 through AH-10. Final statuses come from reconciliation/state records; performance comes only from the seven selected runtime summaries.

6. Model, weight, cache-format, and backend availability

Availability preserves the full ten-row WB-03 matrix and its terminal status classification.

AV-01 — Model, cache, and backend availability

Test	Model	Cache	Backend	Status
AH-01	gemma-1b-q4-k-m	f16	cpu	Passed
AH-02	gemma-1b-q4-k-m	tq3_0	cpu	Passed
AH-03	granite-3b-bf16	f16	cpu	Passed
AH-04	granite-3b-bf16	q8_0	cpu	Passed
AH-05	granite-3b-bf16	tq3_0	cpu	Passed
AH-06	granite-8b-q8-0	f16	cpu	Blocked
AH-07	granite-8b-q8-0	tq3_0	cpu	Blocked
AH-08	granite-3b-bf16	f16	sycl-partial	Passed
AH-09	granite-3b-bf16	tq3_0	sycl-partial	Passed
AH-10	granite-8b-q8-0	tq3_0	sycl-partial	Blocked

7. Complete attempt accounting

Each matrix row has one terminal accounting record; the separate rejected AH-09 attempt remains in the canonical attempt ledger.

AT-01 — Complete terminal attempt accounting

Test	Status	Executed	Backend	Reason
AH-01	Passed	Yes	cpu	Completed with explicitly included evidence
AH-02	Passed	Yes	cpu	Completed with explicitly included evidence
AH-03	Passed	Yes	cpu	Completed with explicitly included evidence
AH-04	Passed	Yes	cpu	Completed with explicitly included evidence
AH-05	Passed	Yes	cpu	Completed with explicitly included evidence
AH-06	Blocked	No	cpu	Safety prerequisite blocked inference before a request; classification resolved with cleanup evidence
AH-07	Blocked	No	cpu	Safety prerequisite blocked inference before a request; classification resolved with cleanup evidence
AH-08	Passed	Yes	sycl-partial	Completed with explicitly included evidence
AH-09	Passed	Yes	sycl-partial	Completed with explicitly included evidence
AH-10	Blocked	No	sycl-partial	Safety prerequisite blocked inference before a request; classification resolved with cleanup evidence

8. Performance results and repetition detail

Reader-facing metrics use stable decimal precision; only explicitly admitted samples contribute to the seven runnable rows.

PF-01 — Performance and resource observations

Test	TTFT ms	Decode tok/s	Peak WS MiB	KV MiB	CPU mean %	GPU mean %	Admission
AH-01	140.39	46.323	929.49	26.00	56.90	0.00	included
AH-02	137.75	37.749	911.06	5.69	58.44	0.00	included
AH-03	274.92	15.872	3663.23	160.00	62.04	0.00	included
AH-04	243.23	15.819	3672.14	170.00	61.97	0.00	included
AH-05	252.57	11.536	3575.03	70.00	62.70	0.00	included
AH-06	Not collected	Not collected	8913.65	320.00	Not collected	Not collected	safety evidence only; excluded from formal statistics

Test	TTFT ms	Decode tok/s	Peak WS MiB	KV MiB	CPU mean %	GPU mean %	Admission
AH-07	Not collected	Not collected	8737.70	140.00	Not collected	Not collected	safety evidence only; excluded from formal statistics
AH-08	309.95	15.184	3130.69	320.00	62.03	16.31	included
AH-09	491.00	12.109	2891.66	70.00	63.14	12.61	included
AH-10	Not collected	Not collected	8969.62	70.00	Not collected	Not collected	safety evidence only; excluded from formal statistics

Missing request-window/OS observations remain literal Not collected; safety evidence is not a formal statistic.

9. Quality methodology and results

The seven runnable rows retain 42 P1-P6 scores under the historical harsh content screen. Calibration is Not collected. This method is not directly comparable with OpenVINO.

QL-01 — Historical prompt-score means

Test	Mean /10	Boundary
AH-01	6.2250	Historical route-specific screen
AH-02	3.6583	Historical route-specific screen
AH-03	6.3917	Historical route-specific screen
AH-04	6.7250	Historical route-specific screen
AH-05	3.2750	Historical route-specific screen
AH-08	6.2750	Historical route-specific screen
AH-09	6.0583	Historical route-specific screen

10. Device/backend use and fallback verification

AH-08 used SYCL OpenCL partial placement. Recovered AH-09 used Level Zero level_zero:0 with CPU-resident TQ3 KV and 1/41 layers offloaded. Vulkan built only as supplementary evidence and is not a source-proven controlled TQ3 route.

11. Failures, blocks, deviations, and recovery attempts

Historical events remain distinct from the final unresolved-failure count and retain their original codes and resolutions.

FL-01 — Historical failures, recoveries, and terminal classifications

ID	Scope	Code	Resolution
AH-F01	AH-B02	DEP	Yes
AH-F02	AH-B02	BF	Yes
AH-F03	AH-B05	REPRO	Yes
AH-F04	AH-B05	REPRO	Yes
AH-F05	AH-B06	REPRO	Yes
AH-F06	AH-B05	REPRO	Yes
AH-F07	AH-B05	REPRO	Yes
AH-F08	AH-06/AH-07	SAFETY	Yes - safety-classified
AH-F09	AH-09	RECOVERY	Yes - recovered with bounded partial placement
AH-F11	AH-10	SAFETY	Yes - safety-classified
AH-F10	AH-04/P5	QUAL	Yes

All eleven historical events remain visible. The formal runtime status has zero unresolved failures, not zero historical attempts.

12. Limitations, uncertainty, robustness checks, and claim boundaries

This historical route did not collect every OS metric for safety-stopped rows; those fields are Not collected, never zero. The evidence supports a conditional research comparator, not production Intel GPU integration or a full-GPU TQ3 claim.

Note: Rejected/superseded summaries are provenance evidence only and never contribute to formal aggregates.

13. Reproduction guidance

Use the ordered commands in reproduction/commands.md to rebuild the route from immutable evidence, export the owned Word PDF, finalize it, and validate checksums.

14. Evidence index and hashes

This reader-facing table lists the bounded key authority set. The complete 742-row audit inventory remains in evidence/evidence-index.csv.

EV-01 — Key authenticated evidence

Evidence ID	Role	Repository-relative path	SHA-256
animehacker-3409c4fff1f57d4b988a	controlled-workbook-markdown	docs/testing/workbooks/text-templates/03_animehacker_TQ3_0_Controlled_Retest_Workbook_v1.md	f4b86b1aaca43745c2a701e38c6571968596ef586e8fd775b69aa4b7983770b0
animehacker-ca0c684bb478b8527653	intended-matrix	experiments/manifests/animehacker-tq3-0/retest-matrix.json	c254679cabac29b2cb2f8c6528cf1fb371161ebfb0a11f7607bd464c3f56484c
animehacker-34ab10f0a8210699f64b	reconciliation	experiments/raw-results/animehacker-tq3-0/2026-07-18/reconciliation.json	02f6dea288f1b704747ec07b0adacd2644990e645eb7fa45ff32fa71098ebc15
animehacker-d0a6c059ef53abfb592a	raw-evidence	experiments/raw-results/animehacker-tq3-0/2026-07-18/runtime/state.json	25bd4e668fe6f443ec338f82d8e6db8cbb5e75dc44a10c200a52ce32003b4cbf
animehacker-79e759f5588cae2c50d5	raw-evidence	experiments/raw-results/animehacker-tq3-0/2026-07-18/runtime-recovery/state.json	6bfd0e30352e587ef1036bdee4ea5fd474cd530c50c882412e3e9cef877d9cea
animehacker-e73b2f1205e6c7f8d4d3	reconciliation	experiments/raw-results/animehacker-tq3-0/2026-07-18/build-cpu/reconciliation-final.json	4347b7dd353001b35bb13d5d1c304e970cae533c10b4102dd918e46731d6c28b
animehacker-9088849e3d9565254d7d	reconciliation	experiments/raw-results/animehacker-tq3-0/2026-07-18/build-sycl/reconciliation.json	b60e97f3855bda1ec0e4d9858e40d7f3c21f669735940aa6788c8c1446b01241
animehacker-44a8d70b058f1c5cac52	reconciliation	experiments/raw-results/animehacker-tq3-0/2026-07-18/build-vulkan/reconciliation.json	57871c024dc5c425692ec436974ac397088a6c43192448719e13b61cbadfddda
animehacker-6517f74c71a0537eb162	quality-evidence	experiments/raw-results/animehacker-tq3-0/2026-07-18/quality/adjudication.json	5123c0fd2016159fb2b48ea2bedfa039af30798112d18b6f52fe914322fe521b
animehacker-c1863347fee0fae7618f	raw-evidence	experiments/raw-results/animehacker-tq3-0/2026-07-18/quality-recovery/quality-adjudications.json	4fb179143e2867023d264c5b5d73cbfe3fa5901c81338a78f5e8d41b3cdc37d9
animehacker-17b6142958ab2916e9ed	quality-register	docs/testing/Quality-Evaluation-Register.csv	4d8e327caf025090b61c917152ba3f656209376e2365b11c40eeaf0c3236797c
animehacker-d25f303e05b5147dc9ca	prompt-contract	experiments/granite_turboquant_intel/prompts/fixd-feasibility-prompt-set-v1.json	9ba512818e81e0ba8da3ddc89cf040dd3b779d1edc41db23d3a896d778de807f
animehacker-4ee219bf5ee960cb094d	quality-rubric	experiments/granite_turboquant_intel/rubrics/quality-rubric-v1.json	a36016f66e02c9e28f0938cf81335dc4b522e9f92b7d9bad3031f90b7ef91d90

15. Revision history

Revision 1.1 records the source-authentication and reader-presentation hardening applied after independent review.

RH-01 — Revision history

Revision	Date	Change
1.0	2026-08-31	Initial unified final-results publication from WB-03 v1.5 evidence
1.1	2026-08-31	Authenticated complete source entities and publication-ready evidence presentation